

R Tutorial - Report

Edison Guevara

9 10 2019

Setting up R Studio

```
library(readr)
library(tidyverse)
library(ggplot2)
```

Predicting braking distance of cars based on speed

Importing dataset

The data of different brands of cars has been considered for this analysis. The dataset used contains 50 observations.

```
cars <- read_csv("~/Ubiquim/Data Analytics Course/Module II/Task 1/R Tutorial Data/cars.csv")
```

```
## Parsed with column specification:
## cols(
##   'name of car' = col_character(),
##   'speed of car' = col_double(),
##   'distance of car' = col_double()
## )
```

Exploring the data

```
summary(cars)
```

```
## name of car      speed of car  distance of car
## Length:50        Min.      : 4.0    Min.      : 2.00
## Class :character  1st Qu.:12.0    1st Qu.: 26.00
## Mode  :character  Median :15.0    Median : 36.00
##                Mean   :15.4    Mean   : 42.98
##                3rd Qu.:19.0    3rd Qu.: 56.00
##                Max.   :25.0    Max.   :120.00
```

```
str(cars)
```

```
## Classes 'spec_tbl_df', 'tbl_df', 'tbl' and 'data.frame': 50 obs. of 3 variables:
## $ name of car : chr "Ford" "Jeep" "Honda" "KIA" ...
## $ speed of car : num 4 4 7 7 8 9 10 10 10 11 ...
```

```
## $ distance of car: num  2 4 10 10 14 16 17 18 20 20 ...
## - attr(*, "spec")=
## .. cols(
## ..   'name of car' = col_character(),
## ..   'speed of car' = col_double(),
## ..   'distance of car' = col_double()
## .. )
```

```
attributes(cars)
```

```
## $names
## [1] "name of car"      "speed of car"      "distance of car"
##
## $class
## [1] "spec_tbl_df" "tbl_df"          "tbl"             "data.frame"
##
## $row.names
## [1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
## [24] 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46
## [47] 47 48 49 50
##
## $spec
## cols(
##   'name of car' = col_character(),
##   'speed of car' = col_double(),
##   'distance of car' = col_double()
## )
```


For convenien the headers of the data frame have been changed as follows:

```
names(cars) <- c("brand", "speed", "distance")
```


Histograms

```
hist(cars$speed, breaks = 25)
```

Histogram of cars\$speed




```
hist(cars$distance, breaks = 120)
```




As can be seen in the histogram of distance, there is one outlier (observation on the far right). This can be more clearly seen in the boxplot below.

Other plots


1 -> cars speed 2 -> cars distance






Looking at the boxplot for the cars distance (2) one outlier can be spotted with a value of 120.

Normal Q-Q Plot for cars speed



In the figure above it can be appreciated that the distribution of car speeds is close to a normal distribution.

Normal Q-Q Plot for cars braking distance




The figure above shows that the distribution of cars braking distances is less normal than the one of cars speeds.

Variable transformation

Considering that the brake distance is as per the laws of physics correlated to the speed to the power of two (2), the independent variable will be transformed as speed^2 .


```
cars_trans <- cars
names(cars_trans) <- c("brand", "speed2", "distance")
cars_trans$speed2 <- cars$speed * cars$speed
ggplot(cars_trans, aes(speed2, distance)) + geom_point()
```

It can be seen in the figure above that the relationship between braking distance and speed^2 is closed to a straight line.

Removing outliers

```
outlier_value <- boxplot(cars_trans$distance, plot = F)$out
outlier_index <- which(cars_trans$distance == outlier_value)
cars_no_outlier <- cars_trans[-outlier_index,]
ggplot(cars_no_outlier, aes(speed2, distance)) + geom_point()
```



Modeling - Linear regression

Creating testing and training sets

```
set.seed(123)
trainsize <- round(nrow(cars_no_outlier)*.7)
testsize <- nrow(cars_no_outlier) - trainsize
trainsize
```

```
## [1] 34
```

```
testsize
```

```
## [1] 15
```

```
training_indices <- sample(seq_len(nrow(cars_no_outlier)), trainsize)
trainset <- cars_no_outlier[training_indices,]
testset <- cars_no_outlier[-training_indices,]
```

Trainset

```
## # A tibble: 34 x 3
```

```
##      brand      speed2 distance
##      <chr>      <dbl>    <dbl>
##  1 GM           289        46
##  2 Nissan        144        26
##  3 Mitsubishi    144        26
##  4 Honda         49         10
##  5 Dodge         400        68
##  6 Acura         400        70
##  7 Mitsubishi    361        56
##  8 Jeep          576        85
##  9 Jeep          225        36
## 10 Honda         225        36
## # ... with 24 more rows
```

Testset

```
## # A tibble: 15 x 3
##      brand      speed2 distance
##      <chr>      <dbl>    <dbl>
##  1 Ford         16         2
##  2 Jeep         16         4
##  3 BMW          81        16
##  4 GMC         169        26
##  5 Chrysler     169        28
##  6 Acura        196        32
##  7 Chevrolet    196        34
##  8 Ford         225        34
##  9 Hyundai     324        48
## 10 Infiniti     324        50
## 11 Mazda       361        54
## 12 Nissan       361        56
## 13 Chrysler    400        66
## 14 Audi        484        76
## 15 Buick       576        84
```

Training the model - linear regression

In the table below the error metrics of the linear regression are summarized. To highlight is the following:

- The R2 is of 0.99 which denotes a very good fit to the data points.
- t-values are 3.982 for Intercept and 55.337 for speed's coefficient. Since both are > 2 we can say there is a strong correlation between the distance and speed, which we expect to be.
- p-value: $< 2.2e-16$. P-value lower than 0.05 accounts also for a high correlation

```
lm_cars <- lm(distance ~ speed2, trainset)
summary(lm_cars)
```

```
##
## Call:
## lm(formula = distance ~ speed2, data = trainset)
##
## Residuals:
```

```
##      Min      1Q Median      3Q      Max
## -4.861 -1.218 -0.597  1.139  6.611
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 3.227165   0.810508   3.982 0.000369 ***
## speed2      0.150405   0.002718  55.337 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.423 on 32 degrees of freedom
## Multiple R-squared:  0.9897, Adjusted R-squared:  0.9893
## F-statistic: 3062 on 1 and 32 DF,  p-value: < 2.2e-16
```

Testing the model - prediction on test sample

```
prediction_cars <- predict(lm_cars, testset)
error_pred_cars <- testset$distance - prediction_cars
testset <- cbind(testset, prediction_cars, error_pred_cars)
testset$abs_error <- abs(testset$error_pred_cars)
```

Testset showing prediction and error of predictions

```
testset
```

```
##      brand speed2 distance prediction_cars error_pred_cars abs_error
## 1      Ford    16        2      5.633647      -3.63364683  3.63364683
## 2      Jeep    16        4      5.633647      -1.63364683  1.63364683
## 3      BMW     81       16     15.409978       0.59002178  0.59002178
## 4      GMC    169       26     28.645627      -2.64562688  2.64562688
## 5  Chrysler   169       28     28.645627      -0.64562688  0.64562688
## 6      Acura   196       32     32.706565      -0.70656454  0.70656454
## 7  Chevrolet   196       34     32.706565       1.29343546  1.29343546
## 8      Ford   225       34     37.068312      -3.06831239  3.06831239
## 9     Hyundai  324       48     51.958417      -3.95841713  3.95841713
## 10 Infiniti   324       50     51.958417      -1.95841713  1.95841713
## 11     Mazda   361       54     57.523406      -3.52340577  3.52340577
## 12     Nissan   361       56     57.523406      -1.52340577  1.52340577
## 13  Chrysler   400       66     63.389205       2.61079539  2.61079539
## 14     Audi   484       76     76.023233      -0.02323287  0.02323287
## 15     Buick   576       84     89.860502      -5.86050192  5.86050192
```


The MAE of the prediction is the following:

```
mean(testset$abs_error)
```

```
## [1] 2.245004
```

In the figure below it can be seen the data points of the testset together with the predicted distance values.

```
ggplot(testset, aes(speed2)) + geom_point(aes(y = distance), color = "blue") + geom_line(aes(y = predic
```



Predicting petal width based on petal length

Importing dataset

```
IrisDataset <- read.csv("~/Ubiquum/Data Analytics Course/Module II/Task 1/R Tutorial Data/iris.csv", head
```

Exploring the data

```
attributes(IrisDataset)
```

```
## $names
## [1] "X"          "Sepal.Length" "Sepal.Width"  "Petal.Length"
## [5] "Petal.Width" "Species"
##
## $class
## [1] "data.frame"
##
## $row.names
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17
## [18] 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34
## [35] 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51
## [52] 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68
## [69] 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85
## [86] 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100 101 102
## [103] 103 104 105 106 107 108 109 110 111 112 113 114 115 116 117 118 119
## [120] 120 121 122 123 124 125 126 127 128 129 130 131 132 133 134 135 136
## [137] 137 138 139 140 141 142 143 144 145 146 147 148 149 150
```

```
summary(IrisDataset)
```

```
##      X      Sepal.Length  Sepal.Width  Petal.Length
## Min.   : 1.00   Min.   :4.300   Min.   :2.000   Min.   :1.000
## 1st Qu.: 38.25   1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600
## Median : 75.50   Median :5.800   Median :3.000   Median :4.350
## Mean   : 75.50   Mean   :5.843   Mean   :3.057   Mean   :3.758
## 3rd Qu.:112.75   3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100
## Max.   :150.00   Max.   :7.900   Max.   :4.400   Max.   :6.900
## Petal.Width      Species
## Min.   :0.100   setosa      :50
## 1st Qu.:0.300   versicolor:50
## Median :1.300   virginica  :50
## Mean   :1.199
## 3rd Qu.:1.800
## Max.   :2.500
```

```
str(IrisDataset)
```


```
## 'data.frame': 150 obs. of 6 variables:
## $ X : int 1 2 3 4 5 6 7 8 9 10 ...
## $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...
```

```
names(IrisDataset)
```

```
## [1] "X" "Sepal.Length" "Sepal.Width" "Petal.Length"
## [5] "Petal.Width" "Species"
```


```
IrisDataset$Species<- as.numeric(IrisDataset$Species)
hist(IrisDataset$Species)
```

Histogram of IrisDataset\$Species




```
hist(IrisDataset$Petal.Length, breaks = 50)
```

Histogram of IrisDataset\$Petal.Length

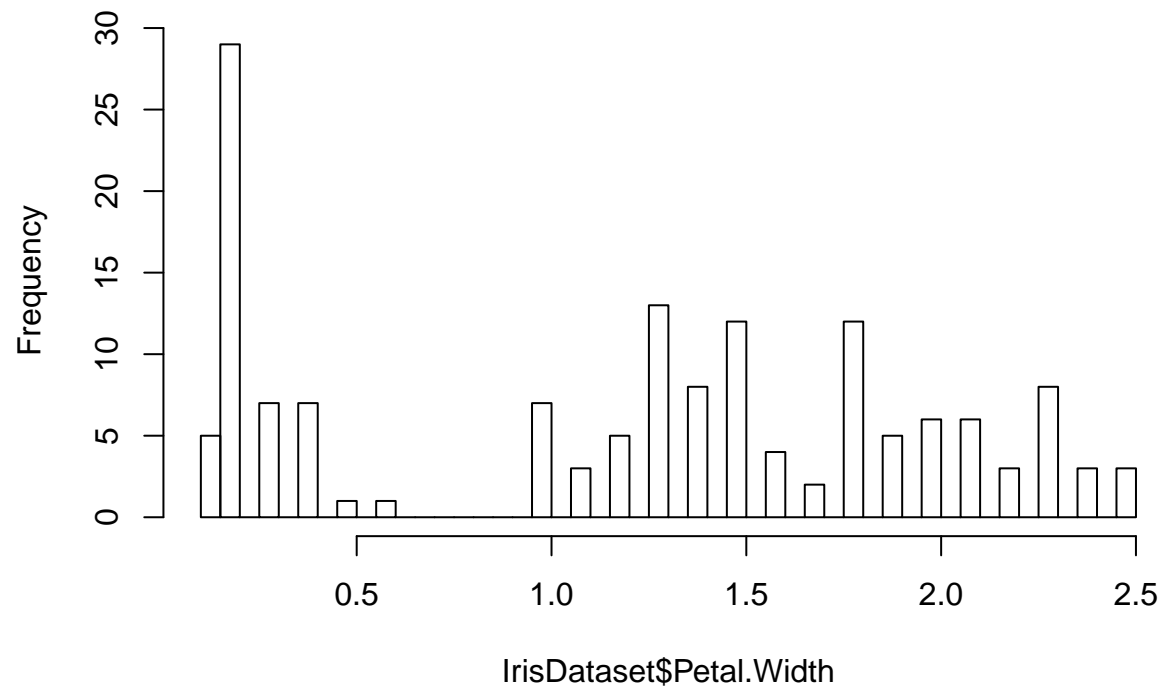


```
boxplot(IrisDataset$Petal.Length)
```





```
hist(IrisDataset$Petal.Width, breaks = 50)
```


Histogram of IrisDataset\$Petal.Width



```
boxplot(IrisDataset$Petal.Width)
```




```
ggplot(IrisDataset, aes(Petal.Length, Petal.Width)) + geom_point()
```



```
qqnorm(IrisDataset$Petal.Length)
```

Normal Q-Q Plot



Training the model - linear regression

```
set.seed(123)
trainSize <- round(nrow(IrisDataset) * 0.2)
testSize <- nrow(IrisDataset) - trainSize
trainSize

## [1] 30

testSize

## [1] 120

train_indic <- sample(seq_len(nrow(IrisDataset)), trainSize)
trainSet <- IrisDataset[train_indic, ]
testSet <- IrisDataset[-train_indic, ]
set.seed(405)
trainSet <- IrisDataset[train_indic, ]
testSet <- IrisDataset[-train_indic, ]
lm_iris <- lm(Petal.Width ~ Petal.Length, trainSet)
summary(lm_iris)
```

```
##
## Call:
## lm(formula = Petal.Width ~ Petal.Length, data = trainSet)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.36964 -0.10766  0.00591  0.08338  0.47607
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -0.28053     0.07150  -3.923 0.000516 ***
## Petal.Length   0.39365     0.01684  23.381 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1742 on 28 degrees of freedom
## Multiple R-squared:  0.9513, Adjusted R-squared:  0.9495
## F-statistic: 546.7 on 1 and 28 DF,  p-value: < 2.2e-16
```

Testing the model - prediction on test sample

```
prediction_iris <- predict(lm_iris, testSet)
error_pred_iris <- testSet$Petal.Width - prediction_iris
testSet <- cbind(testSet, prediction_iris, error_pred_iris)
ggplot(testSet, aes(Petal.Length)) + geom_point(aes(y = Petal.Width), color = "blue") + geom_line(aes(y
```

